

Document explanation and getting started ▾

Share

Good plan — test locally first for speed, batch the Docker rebuild at the end once all three optional features are done. Let's proceed with the OCR code.

Add the `loader.py` changes from my earlier message (the `_needs_ocr` and `_ocr_pdf` functions, plus the updated `load_document_text`).

Now let's create a genuinely scanned-style PDF to test against — since a normal PDF won't trigger the OCR path at all. Easiest way: take a screenshot of some text, then convert it to a PDF with no real text layer.

```
bash
cd sample_data/imtithal
```

Do you have a way to quickly create a scanned-looking PDF — e.g., take a screenshot of a paragraph of text (Cmd+Shift+4 on Mac), save it as a `.png`, and then we'll convert that image into a PDF? That would have zero extractable text layer, which is exactly what triggers your `_needs_ocr` fallback. Let me know once you have a screenshot saved, or if you'd rather I show you a quick way to generate one programmatically instead (e.g., render some text onto a blank image using Python, then save as PDF) — faster if you don't want to fuss with manual screenshots.

just now

Write a message...

 $+$ 

Sonnet 5 Medium ▾

